

Unicode and MultiLingual Computing

Created in Master PDF Editor

Contents

- What is Unicode
- Who uses Unicode
- Benefits of using Unicode
- About the Unicode Consortium
- Encoding basics

What is Unicode

an international encoding standard for use with different languages and scripts, by which each letter, digit, or symbol is assigned a unique numeric value that applies across different platforms and programs. - Source (Definition provided by Google)

*Unicode provides a unique number for every character,
no matter what the platform,
no matter what the program,
no matter what the language.*

- Source (<http://unicode.org/standard/WhatIsUnicode.html>)

Who Uses Unicode

- The Unicode standard has been adopted by such industry leaders as Apple, Google, HP, IBM, JustSystems, Microsoft, Oracle, SAP, Sun, Sybase, Unisys and many others.
- Unicode is required by modern standards such as XML, Java, ECMAScript(JavaScript), LDAP, CORBA 3.0, WML etc. and is the official way to implement ISO/IEC 10646.
- It is supported in many operating systems, all modern browsers, and many other products.
- The emergence of the Unicode Standard, and the availability of tools supporting it, are among the most significant recent global software technology trends.

Benefits of Using Unicode

- Incorporating Unicode into client-server or multi-tiered applications and websites offers significant cost savings over the use of legacy character sets.
- Unicode enables a single software product or a single website to be targeted across multiple platforms, languages and countries without re-engineering.
- It allows data to be transported through many different systems without corruption.

About the Unicode Consortium (www.unicode.org)

- The Unicode Consortium was founded to develop, extend and promote use of the Unicode Standard, which specifies the representation of text in modern software products and standards.
- The Consortium is a non-profit, charitable organization.
- The membership of the Consortium represents a broad spectrum of corporations and organizations in the computer and information processing industry.

Encoding Basics

- Please go through the article - “The Absolute Minimum Every Software Developer Absolutely, Positively Must Know About Unicode and Character Sets (No Excuses!)” by Joel Spolsky, co-founder of Trello and Fog Creek Software, and CEO of Stack Overflow.
- FAQ from the Official Unicode website, http://www.unicode.org/faq//utf_bom.html
- Why do we need Unicode?
 - Lets say Joe Average is a software developer. He insists that he will only ever need English, and as such only wants to use ASCII. This might be fine for Joe the *user*, but this is not fine for Joe the *software developer*. Approximately, half the world uses the non-Latin characters and using ASCII is arguably inconsiderate to these people, and on top of that, he is closing off his software to a large and growing economy.
 - Therefore, an encompassing character set including *all* languages is needed. Thus came Unicode. It assigns every character a unique number called a **code point**.

Encoding Basics

- One advantage of Unicode over other possible sets is that the first 256 code points are identical to ISO-8859-1, https://en.wikipedia.org/wiki/ISO/IEC_8859-1, and hence also ASCII.
- In addition, the vast majority of commonly used characters are represented by only two bytes, in a region called the Basic Multilingual Plane (BMP), https://en.wikipedia.org/wiki/Plane_%28Unicode%29#Basic_Multilingual_Plane
- Memory considerations
 - UTF-8:
 - 1 byte: Standard ASCII
 - 2 bytes: Arabic, Hebrew, most European scripts (most notably excluding Georgian, https://en.wikipedia.org/wiki/Georgian_scripts)
 - 3 bytes: BMP
 - 4 bytes: All Unicode characters
 - UTF-16:
 - 2 bytes: BMP
 - 4 bytes: All Unicode Characters

Encoding Basics

- It's worth mentioning that characters not in the BMP include ancient scripts, mathematical symbols, musical symbols, and rarer Chinese/Japanese/Korean (CJK) characters, http://www.unicode.org/faq/han_cjk.html
- If you'll be working mostly with ASCII characters, then UTF-8 is certainly more memory efficient.
- However, if you are working mostly with non-European scripts, using UTF-8 could be up to 1.5 times less memory efficient than UTF-16.
- When dealing with large amounts of text, such as large web-pages or lengthy word documents, this could impact performance.
- **UTF-8:** For the standard ASCII (0-127) characters, the UTF-8 codes are identical. This makes UTF-8 ideal if backwards compatibility is required with existing ASCII text. Other characters require anywhere from 2-4 bytes. This is done by reserving some bits in each of these bytes to indicate that it is part of a multi-byte character. In particular, the first bit of each byte is 1 to avoid clashing with the ASCII characters.

Encoding Basics

- **UTF-16:** For valid BMP characters, the UTF-16 representation is simply its code point. However, for non-BMP characters UTF-16 introduces surrogate pairs. In this case a combination of two two-byte portions map to a non-BMP character. These two-byte portions come from the BMP numeric range, but are guaranteed by the Unicode standard to be invalid as BMP characters. In addition, since UTF-16 has two bytes as its basic unit, it is affected by endianness. To compensate, a reserved byte order mark can be placed at the beginning of a data stream which indicates endianness. Thus, if you are reading UTF-16 input, and no endianness is specified, you must check for this.
- As can be seen, UTF-8 and UTF-16 are nowhere near compatible with each other. So if you're doing I/O, make sure you know which encoding you are using! For further details on these encodings, please see the UTF FAQ, http://www.unicode.org/faq/utf_bom.html

Encoding Basics

- **Practical programming considerations**
 - **Character and String data types:** How are they encoded in the programming language? If they are raw bytes, the minute you try to output non-ASCII characters, you may run into a few problems. Also, even if the character type is based on a UTF, that doesn't mean the strings are proper UTF. They may allow byte sequences that are illegal. Generally, you'll have to use a library that supports UTF, such as ICU for C, C++ and Java. In any case, if you want to input/output something other than the default encoding, you will have to convert it first.
 - **Recommended/default/dominant encodings:** When given a choice of which UTF to use, it is usually best to follow recommended standards for the environment you are working in. For example, UTF-8 is dominant on the web, and since HTML5, it has been the recommended encoding. Conversely, both .NET and Java environments are founded on a UTF-16 character type. Confusingly (and incorrectly), references are often made to the "Unicode encoding", which usually refers to the dominant UTF encoding in a given environment.

Encoding Basics

- **Practical programming considerations**
 - **Library support:** What encodings are the libraries you are using support? Do they support the corner cases? Since necessity is the mother of invention, UTF-8 libraries will generally support 4-byte characters properly, since 1, 2, and even 3 byte characters can occur frequently. However, not all purported UTF-16 libraries support surrogate pairs properly since they occur very rarely.
 - **Counting characters:** There exist combining characters in Unicode. For example the code point U+006E (n), and U+0303 (a combining tilde) forms ñ, but the code point U+00F1 forms ñ. They should look identical, but a simple counting algorithm will return 2 for the first example, 1 for the latter. This isn't necessarily wrong, but may not be the desired outcome either.
 - **Comparing for equality:** А, А, and Α look the same, but they're Latin, Cyrillic, and Greek respectively. You also have cases like С and С, one is a letter, the other a Roman numeral. In addition, we have the combining characters to consider as well. For more info see Duplicate characters in Unicode, https://en.wikipedia.org/wiki/Duplicate_characters_in_Unicode.

Encoding Basics

- **Practical programming considerations**

- **Surrogate pairs:** These come up often enough on. So, here are some example links

- Getting string length,

<http://stackoverflow.com/questions/12907022/python-getting-correct-string-length-when-it-contains-surrogate-pairs>

- Removing surrogate pairs,

<http://stackoverflow.com/questions/12867000/how-to-remove-surrogate-characters-in-java>

- Palindrome checking

<http://stackoverflow.com/questions/12340930/string-is-palindrome-or-not/12341017#12341017>